

# W05-Arch-Practical-Report

CS2002 170011474 February 25<sup>th</sup>

## Overview

This practical requires students to accomplish three parts: commenting function `itos_recur`, print out all values in the current stack frame of `print_recur()`, analyse the printed result, and analyse the optimization for given assembly codes.

## Implementation & Design (Part 2)

A while loop is used. During every iteration, the offset will be decremented by 8 (bytes). The most critical part is to disassemble each 8 bytes into two 4 bytes and display them separately. Two doublewords are printed firstly and the quadword comes last.

## Reflection of Part 1

1. Discussion about `idiv` and `cqto`:
  - a) `idiv` does a 128/64 bit division, so the value of `%rdx`. In this case, `%rax` holds lower 64 bits from the dividend long val (which is just long val itself), and `%rdx` holds the upper 64 bits. However, `%rdx` must not contain a random value, so the current zero extension is wrong. It needs the sign extension.
  - b) `cqto` (convert quadword to octoword) does a sign extension for `%rdx`: sign extend `%rax` to (`%rdx: %rax`). In this way, the division operation is correct and safe.
2. Conversion between different size types (typically `movslq`):

`movslq` means doing a sign extension when copying 32-bit value into a 64-bit register. This specifies the extension style and avoids inconvenience and unsafety for signed values.

## Analysis of Part 2 (stack frame)

1. Description of values found in the stack frames.

```

=====
@ 140722639540052 | offset: 4 | (doubleword) value is: 32764
@ 140722639540048 | offset: 0 | (doubleword) value is: -1963913312
@ 140722639540048 | offset: 0 | (quadword) value is: 140722639540128
-----
@ 140722639540044 | offset: -4 | (doubleword) value is: -1
@ 140722639540040 | offset: -8 | (doubleword) value is: -1
@ 140722639540040 | offset: -8 | (quadword) value is: -1
-----
@ 140722639540036 | offset: -12 | (doubleword) value is: 2
@ 140722639540032 | offset: -16 | (doubleword) value is: 0
@ 140722639540032 | offset: -16 | (quadword) value is: 8589934592
-----
@ 140722639540028 | offset: -20 | (doubleword) value is: 32764
@ 140722639540024 | offset: -24 | (doubleword) value is: -1963913208
@ 140722639540024 | offset: -24 | (quadword) value is: 140722639540232
-----
@ 140722639540020 | offset: -28 | (doubleword) value is: 0
@ 140722639540016 | offset: -32 | (doubleword) value is: 0
@ 140722639540016 | offset: -32 | (quadword) value is: 0
-----
@ 140722639540012 | offset: -36 | (doubleword) value is: 0
@ 140722639540008 | offset: -40 | (doubleword) value is: 1
@ 140722639540008 | offset: -40 | (quadword) value is: 1
-----
@ 140722639540004 | offset: -44 | (doubleword) value is: 0
@ 140722639540000 | offset: -48 | (doubleword) value is: 8229472
@ 140722639540000 | offset: -48 | (quadword) value is: 8229472
-----
@ 140722639539996 | offset: -52 | (doubleword) value is: 5
@ 140722639539992 | offset: -56 | (doubleword) value is: -350628112
@ 140722639539992 | offset: -56 | (quadword) value is: 25419175664
-----
@ 140722639539988 | offset: -60 | (doubleword) value is: 32764
@ 140722639539984 | offset: -64 | (doubleword) value is: -1963913088
@ 140722639539984 | offset: -64 | (quadword) value is: 140722639540352
-----
@ 140722639539980 | offset: -68 | (doubleword) value is: 0
@ 140722639539976 | offset: -72 | (doubleword) value is: 4199168
@ 140722639539976 | offset: -72 | (quadword) value is: 4199168
=====
133

```

offset	value	Purpose
8	return address (8 bytes)	return to the code segment to the next instruction in the calling method.
0	the address of the previous %rbp (8 bytes)	the control flow will jump back to the previous function after terminating.
-8	long val (8 bytes)	the first argument
-12	int depth (4 bytes)	the second argument
-16	0	16-byte alignment
-24	char**end (8 bytes)	the third argument; the address of a pointer to a char array.
-32	long quotient (8 bytes)	
-40	long mod (8 bytes)	
-48	return value in %rax (8 bytes)	
-52	int len (4 bytes)	length of the string
-56, -64	0	16-byte alignment
-72	return address (8 bytes) as above	

2. **Tests:** Results are correct and correspond to my analysis above. Human readers can access values from the “dump” by those straightforward doublewords and quadwords.

### 3. Utility of the stack frames and explanation

A stack frame under this situation consists of **64 bytes** (or **80** regardless of CFA), since ***subq \$64, %rsp*** indicates 64 bytes reserved. Between two frames there are **16** bytes holding *return address* and the pushed *previous %rbp*. This convention is defined by *CFA (Canonical Frame Address)*: it is the value of the stack pointer at the call site in the previous frame.

There are 16 bytes not used, 4 bytes following *int depth* and 12 bytes following *%rax*.

**The reason is that CFA convention has set the offset 16 for alignment.** Every block of 16 bytes in stack frame is seen as a unit to store values, starting from *%rbp*. For example, *long val* occupies the first 8 bytes, *int depth* occupies the next 4 bytes. Since there is only 4 bytes left in the current 16-byte unit, which means the following *char \*\*end* cannot be aligned in the same unit. For maximum loading efficiency, *char \*\*end* is stored in the beginning of the next unit (-17 to -32), **leaving those 4 bytes empty 0.**

4. x86 adopts **little-endian** to store values. If we focus on the offset -8 with value 123 stored, from -1 to -4 there is a 0, and from -5 to -8 there is a 123. Thus, x86 stores less significant values in less significant bits.
5. The effect brought by added function `print_stack()` on format of stack frames

This is no dramatic change in the layout of stack frames. Although recursively calling `print_recur()` pushes new frames onto the stack, positions of values stored in the current frame are not changed.

## Analysis of Part 3 (Division by invariant integers using multiplication)

### 1. Error Analysis of Formula (2) and Formula (3):

The formula (2) does not work for division. For  $n = 2^{63} - 1$ ,  $2^{65} \equiv 2 \pmod{10}$ , hence  $\lceil 2^{65}/10 \rceil - 2^{65}/10 = 8/10$ . Moreover,  $(n/10 - \lfloor n/10 \rfloor)$  can be up to  $9/10$ . The upper bound of the error is  $8/10 * (2^{63} - 1)/2^{65} + 9/10 \approx 1.1 \geq 1$ . This shows formula (2) has the same problem as formula (1).

The formula (3) gives us desired result.  $2^{66} \equiv 4 \pmod{10}$ , hence  $\lceil 2^{66}/10 \rceil - 2^{66}/10 = 6/10$ . The upper bound of the error is  $6/10 * (2^{63} - 1)/2^{66} + 9/10 \approx 0.975 < 1$ . Since the error is less than 1, it can be safely ignored.

### 2. Implementation of $2^{66}$ in the assembly file (5 lines) for non-negative n:

- a) `movq %rdi, %r15`
- b) `movabsq $7378697629483820647, %rcx`
- c) `movq %r15, %rax`
- d) `imulq %rcx`
- e) `sarq $2, %rdx`

The argument *long val* was copied into *%rax* after (a), (c). We can notice that 7378697629483820647 is equal to  $2^{66}/10 = \lceil 2^{66}/10 \rceil$  in the formula. (b) is to copy this number into *%rcx*. In (d), the upper 64 bits of the 128-bit product is put into *%rdx* while

the lower 64 bits is put into `%rax`. In (e), the upper 64 bits right shift by 2 bits, which indicates now the remaining number in `%rdx` is  $(n/2^{66})$ . This is what we are looking for.

### 3. Handling negative n:

There are 2 fundamental operations: obtain quotients and mods.

- a) For long quot, we can obtain a quot after the multiplication and right shift. However, under signed division and quotient rounded towards 0, if  $n < 0$  while the divisor is positive then the current quotient should be 1 less than the true quotient (Granlund and Montgomery, 1994). To recover the correct value, the quotient is incremented by 1. What has to be mentioned is that this 1 can come from the sign bit of long val. If  $n \geq 0$ , then the quotient remains the same; If  $n < 0$ , then the sign bit is 1 which will be added to the quotient.
- b) To obtain an absolute mod, the mod is negated by using `negq`.

## Extension

### 1. Further optimisations in `itos1.s`

- a) When a conditional statement comes, if `(quot != 0)`, it uses 19 compared to `val + 9` to see if `val` is greater than 10, which can be used to determine the quotient.
- b) `leaq` is a very quick way to calculate a result directly stored into the destination register. In `itos1.s`, `leaq` is adopted in many arithmetic calculations, which can reduce the compile time.
- c) Before defining the int len, load the sign bit from long val into `%rax`. If long val is negative, then the sign bit is 1, otherwise the sign bit is 0. Afterwards, int len is incremented by 2. This avoids a comparison between val and 0 and reduces the number of registers to be used.
- d) When a comparison involves 0 or setting a register with 0, `testq` and `xorl` can achieve them in a logical bitwise way. This is quick and safe within just 1 instruction.

### 2. Comparison among MIPS, X86, ARM

Differences	MIPS	ARM	x86
CISC / RISC	RISC	RISC	CISC
Destination registers	Usually the first operand	Usually the first operand	Usually the last operand
Locate the specific value	Offset and <code>\$sp</code>	offset and <code>sp</code> , <code>r11(fp)</code>	Offset and <code>rbp</code>

- a) CISC processors tend to use a single instruction with multiple tasks finished. For example, when operating an expression like  $(x + 4*y)$ , it can use `leaq (%reg1, %reg2, 4)`, or when doing a comparison, the flags have also been set. However, RISC processors are likely to implement one task by one instruction.
- b) When a branch instruction appears, MIPS will use "nop" to fill the branch delay slot of jumps or branches in the pipeline. Since if the pipeline is too "deep", a branch will lead to many wasteful instructions following the branch, and they should not be used after branching to another place. "nop" can take up 5 stages

in pipeline and it can reduce the cost in "branching time". X86 so far in itos1.s has no similar method.

- c) **Calling conventions:** MIPS usually calls functions with the help of jal and jr \$ra explicitly. ARM has similar techniques to *"jump and link"* and *jump back* at the function epilogue. This difference from x86 leads to extra operations to push and pop these registers. Additionally, MIPS and ARM have 2 general registers to hold returning values, while x86 only has *%rax*. Moreover, MIPS contains 2 special registers called *\$hi* and *\$lo*. They hold values from multiplication and division, while x86 uses *%rdx* and *%rax* but which can also be used to store other values.

## Evaluation & Conclusion

My work satisfies requirements and my report has answered all questions raised in the specification. Additionally, in order to understand the architecture of stack frames extensively, I also searched many articles to consolidate my foundation.

In conclusion, I have acquired the deeper insights into the architecture of stack frames and assembly codes of X86. I still need to learn more from MIPS.

## Reference list

Granlund, T. and Montgomery, P. (1994). Division by invariant integers using multiplication. *ACM SIGPLAN Notices*, [online] 29(6), pp.61-72. Available at: <https://gmplib.org/~tege/divcnst-pldi94.pdf> [Accessed 24 Feb. 2019].